

# Reinforcement Learning Visualizer - Technical Report

## 1. Project Overview

The **Reinforcement Learning Visualizer** is an interactive web application built to demonstrate and visualize core concepts of Reinforcement Learning (RL). It provides a hands-on learning experience, allowing users to experiment with different RL algorithms, environments, and parameters to better understand how reinforcement learning works in practice.

## 2. Implemented Environments

### 2.1 GridWorld (5×5 Grid)

- **Type:** Discrete state space
- **Actions:** 4 (Up, Right, Down, Left)
- **Features:**
  - Goal state located at the bottom-right corner
  - Obstacles placed at positions [6, 12]
  - Deterministic transitions (fixed outcomes)
  - Complete transition model available
- **Rewards:** Users can configure the goal reward, step penalty, and obstacle penalty.

### 2.2 Frozen Lake (4×4 Grid)

- **Type:** Discrete state space
- **Actions:** 4 (Up, Right, Down, Left)
- **Features:**
  - Slippery transitions (33% chance of slipping)
  - Holes placed at positions [5, 7, 11, 12]
  - Goal at the bottom-right corner
  - Complete transition model available
- **Rewards:** Users can adjust the goal reward, hole penalty, and step penalty.

### 2.3 CartPole

- **Type:** Continuous state space (discretized into 100 states)
- **Actions:** 2 (Left, Right)
- **Features:**

- Simulates continuous physics (pole balancing)
- State discretization based on position (10 bins) and angle (10 bins)
- No transition model required
- **Rewards:** Customizable step reward and fail penalty.

## 2.4 Mountain Car

- **Type:** Continuous state space (discretized into 100 states)
- **Actions:** 3 (Left, Nothing, Right)
- **Features:**
  - Simulates continuous physics, including gravity and friction
  - State discretization based on position (10 bins) and velocity (10 bins)
  - Goal state at position  $\geq 0.5$
  - No transition model required
- **Rewards:** Users can adjust goal reward and step penalty.

## 2.5 Breakout Environment

- **Type:** Discrete state space (200 states)
- **Actions:** 3 (Left, Stay, Right)
- **Features:**
  - A simplified version of the classic Breakout game
  - Real-time physics for paddle, ball, and bricks
  - Real-time collision detection
  - No transition model required
- **Rewards:** Configurable rewards for brick destruction, miss penalties, and win conditions.

## 3. Implemented Algorithms

### 3.1 Dynamic Programming Methods

#### Policy Evaluation

- **Category:** DP (Requires transition model)
- **Description:** Evaluates the value function of a given policy.
- **Use Case:** GridWorld, Frozen Lake

#### Policy Iteration

- **Category:** DP (Requires transition model)
- **Description:** Alternates between policy evaluation and improvement.
- **Use Case:** GridWorld, Frozen Lake

#### Value Iteration

- **Category:** DP (Requires transition model)
- **Description:** Computes the optimal value function directly.
- **Use Case:** GridWorld, Frozen Lake

## 3.2 Model-Free Methods

### Monte Carlo

- **Category:** Model-Free (No transition model needed)
- **Description:** Learns from complete episodes.
- **Use Case:** All environments

### TD(0)

- **Category:** Temporal Difference
- **Description:** One-step temporal difference learning.
- **Use Case:** All environments

### n-step TD

- **Category:** Temporal Difference
- **Description:** Multi-step temporal difference.
- **Use Case:** All environments

## 3.3 Control Algorithms

### SARSA

- **Category:** On-Policy TD Control
- **Description:** On-policy action-value learning.
- **Use Case:** All environments

### Q-Learning

- **Category:** Off-Policy TD Control
- **Description:** Off-policy optimal action-value learning.
- **Use Case:** All environments

## 4. Parameter Adjustment Capabilities

### 4.1 Algorithm Parameters

Parameter	Range	Description	Applicable Algorithms
$\gamma$ (Gamma)	0-1	Discount factor for future rewards	All algorithms
$\alpha$ (Alpha)	0.01-1	Learning rate for updates	Q-Learning, SARSA, TD methods
$\epsilon$ (Epsilon)	0-1	Learning rate for updates	Q-Learning, SARSA, Monte Carlo
n-step	1-10	Look-ahead steps	n-step TD
Iterations	1-500	DP algorithm iterations	Policy/Value Iteration

## 4.2 Training Parameters

- **Episodes:** 10-1000 training episodes
- **Batch Size:** Automatic batch processing for smooth UI updates
- **Progress Tracking:** Real-time progress bar and episode count

## 4.3 Environment Reward Customization

Each environment has customizable reward settings:

- **GridWorld:** Goal reward, step penalty, obstacle penalty
- **Frozen Lake:** Goal reward, hole penalty, step penalty
- **CartPole:** Step reward, fail penalty, max steps
- **Mountain Car:** Goal reward, step penalty, max steps
- **Breakout:** Brick reward, miss penalty, step penalty, win reward

# 5. Visualization Techniques

## 5.1 Environment Visualization

### Grid-based Environments (GridWorld, Frozen Lake)

- **Heat Map:** State values color-coded using a blue gradient.
- **Policy Arrows:** Indicate optimal action directions.
- **Agent Position:** Orange circle representing the agent's position.
- **Special States:**
  - Green for the goal state.
  - Red for obstacles/holes.
- **Interactive Cells:** Click to inspect state values.

### Continuous Environments (CartPole, Mountain Car)

- **Physics Simulation:** Real-time motion visualization.
- **Parameter Displays:** Displays real-time angle, position, and velocity.
- **Force Indicators:** Displays velocity vectors and applied forces.
- **Boundary Markers:** Failure boundaries shown on the environment.

### Breakout Environment

- **Game Elements:** Bricks, paddle, and ball with 3D effects.

- **Score Display:** Real-time score and brick count.
- **Velocity Vector:** Ball direction indicator.
- **Status Overlay:** "PLAYING" indicator during inference.

## 5.2 Training Visualization

### Value Function Display

- **Grid Layout:** Displays state values in matrix format.
- **Color Coding:** Blue intensity represents state values.
- **Tooltips:** Hover to see exact values.
- **Cell Selection:** Click to highlight specific states.

### Metrics Charts

- **Episode Returns:** Green line chart showing reward progression.
- **Convergence ( $\Delta V$ ):** Red line chart showing DP algorithm convergence.
- **Statistics Display:**
  - Last return value
  - Average return
  - Episode count
  - Total steps

## 5.3 Control Panel Visualization

### Status Indicators

- **Badge System:** Color-coded status (Training, Trained, Running, Paused).
- **Parameter Display:** Displays current values for  $\gamma$ ,  $\alpha$ ,  $\epsilon$ .
- **State Counter:** Shows the current environment state.

### Progress Tracking

- **Progress Bar:** Displays training completion percentage.
- **Episode Counter:** Shows the current/total episodes.
- **Real-time Updates:** Live updates of parameter changes.

## 5.4 Interactive Features

- **Real-time Parameter Adjustment:** Sliders provide instant feedback.
- **Environment Switching:** Seamless transitions between environments.
- **Algorithm Comparison:** Compare algorithms side-by-side.
- **Reward Customization:** Dynamically adjust reward settings.
- **Training Control:** Includes start, pause, resume, and reset functions.

## 6. Key Features

1. **Modular Design:** Separate environments, algorithms, and UI components.
2. **Real-time Updates:** Canvas animations and live parameter feedback.
3. **Error Handling:** Graceful handling of NaN/Infinity values.
4. **Responsive Design:** Adapts to different screen sizes.
5. **Performance Optimization:** Batch training and efficient rendering.